

Regime-Adaptive Portfolio Optimization via Forward-Filtered Hidden Markov Models

Drawdown Control Through Regime-Conditional Convex Optimization

Abdelkrim Abdelkrim

Applied Mathematics & Artificial Intelligence

Université Paris Dauphine – PSL

github.com/AbdelkrimCode/regime-adaptive-portfolio

July 22, 2026

Abstract

We present a regime-adaptive portfolio framework that combines a Gaussian Hidden Markov Model (HMM) for market regime detection with regime-conditional convex optimization. The key methodological contribution is the replacement of standard Viterbi decoding and forward-backward smoothing with a strictly causal forward filter – meaning the filter at time t uses only observations through t , eliminating within-window lookahead bias present in most HMM-based portfolio systems. Applied to eight liquid ETFs over 2006–2024, the strategy achieves a maximum drawdown of -25.6% versus -59.6% for SPY buy-and-hold, with a full-period Sharpe ratio of 0.561 versus 0.435. A controlled comparison holding the fitted models fixed shows the causal forward filter trailing forward-backward smoothed posteriors by roughly 0.02 Sharpe – the expected direction, since smoothing benefits from hindsight unavailable in live trading. A frozen out-of-sample evaluation – training exclusively on pre-2019 data – yields a Sharpe of 0.339, ahead of the walk-forward estimate of 0.261 on the same period, confirming performance does not depend on lookahead in the walk-forward procedure. In the held-out 2019–2024 period, the strategy’s Sharpe trails SPY (0.677) by a wide margin, consistent with a sustained equity bull market unfavorable to defensive allocation, and compounded by a corrected risk-parity optimizer (Section 2) that trades some resilience to 2022’s correlated bond/equity selloff for real risk balance. Bootstrap significance testing yields $p = 0.344$; the strategy’s primary defensible value proposition is tail risk control rather than return generation.

1 Introduction

Financial markets exhibit distinct statistical regimes — periods of low volatility trending growth, high volatility mean reversion, and acute crisis episodes — that are not well captured by time-invariant models. Portfolio strategies optimized under a single distributional assumption perform poorly when the underlying regime shifts, as evidenced by the catastrophic drawdowns of traditionally diversified portfolios during the 2008 Global Financial Crisis (GFC) and the 2022 simultaneous equity and bond selloff.

Hidden Markov Models [Hamilton, 1989] provide a principled framework for modeling latent market states. In the portfolio management literature, HMM-based regime detection has been applied to asset allocation [Ang and Bekaert, 2002], volatility forecasting [Turner et al., 1989], and risk management [Guidolin and Timmermann, 2007]. However, most implementations rely on the Viterbi algorithm or forward-backward smoothing for state inference — both of which use future observations within the estimation window, introducing a subtle but consequential form of lookahead bias.

This paper uses “causal” in the signal-processing sense: the regime filter at time t conditions only on observations through t , not future observations. This is distinct from causal inference (do-calculus, interventional distributions); the contribution is methodological rather than causal-theoretic.

This paper makes the following contributions:

1. **Forward-only filter:** We replace Viterbi decoding and forward-backward smoothing with a strictly causal forward filter that computes $P(\text{state}_t \mid \mathbf{o}_{1:t})$ using only past observations, eliminating within-window lookahead bias. Holding the fitted HMM models fixed and varying only the decoding method, the causal forward filter trails smoothed posteriors by roughly 0.02 Sharpe (Section 5). This is the expected direction of effect: smoothed posteriors have access to future observations within each estimation window and so can assign better-fitting labels in hindsight than any live system could achieve. The forward-filter number is the honest one — most published HMM-based portfolio backtests do not report this gap at all, and many implicitly rely on smoothing or Viterbi decoding without acknowledging the resulting lookahead.
2. **Posterior probability blending:** Rather than hard regime assignment, we blend optimizer outputs across all regimes weighted by filtered posterior probabilities, reducing turnover at transition boundaries.
3. **Nested cross-validation:** State count selection via BIC is performed independently at each quarterly fold, preventing information leakage from the full sample into model selection.
4. **Comprehensive validation:** We provide a frozen out-of-sample evaluation, feature ablation study, walk-forward leakage audit, bootstrap significance test, and stress tests — a level of methodological scrutiny absent from most published backtests.

The remainder of this paper is organized as follows. Section 2 describes the full methodology. Section 3 details the data and implementation. Section 4 presents performance results. Section 5 reports robustness analyses. Section 6 discusses limitations, and Section 7 concludes.

2 Methodology

2.1 Feature Construction

We construct three features from daily asset returns:

$$r_t^{\text{SPY}} = \log \left(\frac{P_t^{\text{SPY}}}{P_{t-1}^{\text{SPY}}} \right) \quad (1)$$

$$\sigma_t = \sqrt{\frac{252}{W_\sigma} \sum_{i=0}^{W_\sigma-1} (r_{t-i}^{\text{SPY}})^2} \quad (2)$$

$$\bar{\rho}_t = \frac{1}{N(N-1)} \sum_{i \neq j} \hat{\rho}_{ij,t} \quad (3)$$

where $W_\sigma = 21$ trading days, $\hat{\rho}_{ij,t}$ is the rolling $W_\rho = 63$ -day pairwise correlation between assets i and j , and $N = 8$ is the number of assets. Both window parameters are centralized in `config.yaml` and can be overridden via the sensitivity sweep. Features are standardized to zero mean and unit variance before HMM fitting.

2.2 Gaussian Hidden Markov Model

We model the feature sequence $\{\mathbf{x}_t\}$ as observations from a K -state Gaussian HMM with full covariance matrices:

$$P(\mathbf{x}_t \mid s_t = k) = \mathcal{N}(\mathbf{x}_t; \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \quad (4)$$

$$P(s_t = j \mid s_{t-1} = i) = A_{ij} \quad (5)$$

where $s_t \in \{1, \dots, K\}$ is the latent regime, $\boldsymbol{\mu}_k$ and $\boldsymbol{\Sigma}_k$ are regime-specific mean and full covariance, and A is the $K \times K$ transition matrix. Parameters $\{\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k, A, \boldsymbol{\pi}\}$ are estimated by the Baum-Welch algorithm (EM). Multiple random initializations ($n_{\text{init}} = 10$) are used to mitigate local optima; the initialization with the highest log-likelihood is retained, and each fold's fit is checked against the actual convergence criterion (the log-likelihood improvement dropping below tolerance, as opposed to merely exhausting the iteration budget) rather than assumed. Across the current 75-fold walk-forward run, all 75 folds converged under this check.

2.3 State Count Selection via BIC

The number of states K^* is selected at each quarterly fold by minimizing the Bayesian Information Criterion:

$$\text{BIC}(K) = -2\ell(\hat{\theta}_K) + p_K \log T \quad (6)$$

where $\ell(\hat{\theta}_K)$ is the maximized log-likelihood, T is the number of training observations, and p_K counts the free parameters assuming full covariance matrices:

$$p_K = \underbrace{K(K-1)}_{\text{transition}} + \underbrace{Kd}_{\text{means}} + \underbrace{K \frac{d(d+1)}{2}}_{\text{covariances}} + \underbrace{(K-1)}_{\text{initial state}} \quad (7)$$

with $d = 3$ features. For $K = 4$: $p_4 = 12 + 12 + 24 + 3 = 51$. Candidates $K \in \{2, 3, 4\}$ are evaluated independently at each fold, and the winning fit is reused directly rather than refit from scratch.

2.4 Causal Forward Filter

Standard HMM inference tools compute smoothed posteriors $P(s_t \mid \mathbf{x}_{1:T})$ using forward-backward passes over the full observation sequence. This introduces within-window lookahead: the regime label assigned to day t depends on observations $\mathbf{x}_{t+1}, \dots, \mathbf{x}_T$, which are unavailable in real-time.

We replace smoothing with a forward-only filter. The filtered posterior $\alpha_t(k) = P(s_t = k \mid \mathbf{x}_{1:t})$ is computed recursively:

$$\alpha_1(k) = \pi_k \cdot p(\mathbf{x}_1 \mid s_1 = k) \quad (8)$$

$$\alpha_t(k) = p(\mathbf{x}_t \mid s_t = k) \sum_{j=1}^K \alpha_{t-1}(j) A_{jk} \quad (9)$$

with normalization $\alpha_t \leftarrow \alpha_t / \sum_k \alpha_t(k)$ at each step to prevent underflow. The regime assignment on day t depends exclusively on $\mathbf{x}_{1:t}$. The initial state α_1 is computed using π learned on the training window; in test folds beginning mid-cycle, this initialization introduces a minor bias that cannot be corrected without knowing the true state at the fold start.

2.5 Walk-Forward Retraining

To prevent the HMM from using future market structure in regime detection, we employ quarterly walk-forward retraining with an expanding training window:

1. At each quarterly retraining date τ_q , fit the HMM on all data $\mathbf{x}_{1:\tau_q}$
2. Apply the fitted model via forward filter to observations in the subsequent quarter $[\tau_q + 1, \tau_{q+1}]$
3. Advance to the next quarter and repeat

Train and test windows are non-overlapping by construction. Data collection begins January 2005; the first 252 trading days are used to initialize rolling features and satisfy the minimum training window, so the backtest period begins January 2006. Of the 78 theoretical quarterly retraining dates spanning the data collection window, 3 are excluded because the expanding window does not yet meet the 252-day minimum training requirement, yielding 75 usable folds. A programmatic audit confirms zero leakage across all 75 folds (Section 5).

2.6 Label Permutation Handling

HMM states have no intrinsic semantic identity — state index 0 in quarter q need not correspond to the same market condition as state index 0 in quarter $q + 1$. We resolve this by assigning regime labels based on a risk-adjusted score: each state is ranked by $\text{mean_return} - 0.5 \times \text{mean_vol}$, with the lowest-ranked state labeled Crash, followed by Bear, Sideways, and Bull. The two inputs to this score are resolved by feature column name rather than position, so the ranking remains correct even when the feature set itself changes shape (Section 5 discusses a feature-ablation variant with no volatility column at all, where an earlier

position-based implementation silently substituted a different feature). This risk-adjusted ranking prevents mislabeling of high-volatility states with similar returns as safer regimes. We verify in Section 5 that regime rank ordering is stable across all quarterly folds, with no large drawdown coinciding with a label reassignment event.

2.7 Regime-Conditional Optimization

Four convex optimization programs are defined, one per regime. All use the Ledoit-Wolf shrinkage estimator [Ledoit and Wolf, 2004] for $\hat{\Sigma}$, applied to an expanding window of all returns through the retraining date.

Bull — Maximum Sharpe: Using the Lagrangian auxiliary variable reformulation [Cornuéjols and Tütüncü, 2006]:

$$\min_{\mathbf{y}} \quad \mathbf{y}^\top \hat{\Sigma} \mathbf{y} \quad (10)$$

$$\text{s.t.} \quad (\hat{\boldsymbol{\mu}} - r_f)^\top \mathbf{y} = 1, \quad \mathbf{0} \leq \mathbf{y} \leq 0.60 \cdot \mathbf{1}^\top \mathbf{y} \quad (11)$$

where $\mathbf{w} = \mathbf{y}/\mathbf{1}^\top \mathbf{y}$ and r_f is the time-varying 3-month T-bill rate ($\text{\textasciitilde{IRX}}$). The position-limit constraint is expressed in the transformed variable $\mathbf{y}$ so that, after the change of variables back to $\mathbf{w}$, it correctly enforces $w_i \leq 60\%$ regardless of the (a priori unknown) scale of $\mathbf{y}$.

Bear — Constrained Risk Parity: Weights approximately equalize each asset’s contribution to portfolio variance using the Spinu (2013) log-barrier surrogate [Spinu, 2013]:

$$\min_{\mathbf{w}} \quad \mathbf{w}^\top \hat{\Sigma} \mathbf{w} - \frac{1}{N} \sum_{i=1}^N \log w_i, \quad w_i \geq 0.01 \quad (12)$$

followed by normalization $\mathbf{w} \leftarrow \mathbf{w}/\mathbf{1}^\top \mathbf{w}$, then a cap-and-redistribute step enforcing $w_i \leq 0.60$: any asset exceeding the cap is set to exactly 0.60 and the excess is redistributed proportionally among the remaining assets, iterating until no asset exceeds the limit.

Getting this ordering right took two attempts, and both are reported here because the gap between them is informative. The budget constraint $\mathbf{1}^\top \mathbf{w} = 1$ must be applied *after* solving, not as a constraint inside the optimization problem: the log-barrier objective’s first-order conditions pin down each asset’s risk contribution up to a scale factor, and that property is scale-invariant under post-hoc normalization, but is perturbed by a Lagrange multiplier if the budget constraint sits inside the same problem — an initial version did exactly this, and on a controlled synthetic test the resulting risk-contribution spread widened from approximately 0.002 percentage points (bounds-only, normalized after) to approximately 18 percentage points (equality-constrained inside the optimization). Removing the equality constraint and validating against synthetic i.i.d. returns appeared to resolve this — but that synthetic data never exercised the position-limit cap, since i.i.d. returns keep the optimizer’s natural unconstrained solution scale comfortably under 1.0 per asset. Against real historical returns for the actual 8-asset universe, the position cap $w_i \leq 0.60$ — still applied to the *raw*, *pre-normalization* variable in that version — caused most or all assets to saturate at the cap simultaneously (confirmed on three real historical windows: 14–23 percentage points of risk-contribution spread, with the raw pre-normalization solution pinned at exactly 0.60 for every asset in two of the three windows), since that raw variable’s scale is set by $\kappa = 1/N$ against the covariance matrix’s real magnitude and bears no relationship to summing to 1.

Normalizing N identical capped values simply reproduces equal weight. The final formulation above — floor only inside the optimization, cap applied to the already-normalized weights — resolves this (0.0005–0.0014 percentage points of spread on the same three real windows). The lower bound $w_i \geq 0.01$ keeps the log-barrier term finite. The static Risk Parity benchmark uses the same corrected formulation for comparability.

Sideways — Minimum Variance:

$$\min_{\mathbf{w}} \quad \mathbf{w}^\top \hat{\Sigma} \mathbf{w} \quad (13)$$

$$\text{s.t.} \quad \mathbf{1}^\top \mathbf{w} = 1, \quad 0 \leq w_i \leq 0.60 \quad (14)$$

Position limits: Mean-Variance and Minimum-Variance both enforce $w_i \leq 60\%$ per asset directly inside their respective convex problems — this is valid for both, since neither has the scale-ambiguity property of the log-barrier risk-parity formulation above. If any regime’s final weight vector nonetheless concentrates more than 99% of weight in a single asset — a numerical degeneracy rather than an intended allocation — it is replaced with an equal-weight fallback. For the Crash regime specifically, that fallback is scoped to the regime’s own safe-haven subset $\mathcal{S}$ rather than the full 8-asset universe, since a full-universe fallback would inject equity exposure at precisely the moment the Crash regime exists to avoid it.

Crash — Inverse-Volatility Safe Haven: Weights inversely proportional to trailing volatility over the safe-haven subset $\mathcal{S} = \{\text{IEF}, \text{TLT}, \text{GLD}\}$:

$$w_i = \frac{\sigma_i^{-1}}{\sum_{j \in \mathcal{S}} \sigma_j^{-1}}, \quad i \in \mathcal{S} \quad (15)$$

where σ_i is the trailing volatility of asset i . Since IEF (7–10 year Treasuries) is typically the least volatile of the three safe havens, it typically receives the largest allocation of the three, with TLT and GLD receiving smaller allocations; the exact split varies with realized volatility at each retraining date. All optimizers use an expanding window of returns through the retraining date; early folds are dominated by 2008 volatility while late folds blend multiple regimes (see Section 6).

2.8 Posterior Probability Blending

Rather than committing to a single regime optimizer, daily weights are computed as a convex combination weighted by filtered posteriors:

$$\mathbf{w}_t = \sum_{k \in \mathcal{K}} \alpha_t(k) \cdot \mathbf{w}_t^{(k)} \quad (16)$$

where $\mathbf{w}_t^{(k)}$ is the weight vector produced by regime k ’s optimizer using all returns available through day t . Blending reduces turnover at transition boundaries and propagates the HMM’s regime uncertainty into the portfolio allocation. All four optimizer weight vectors are recomputed together whenever the dominant regime changes or at each quarterly retraining date, ensuring consistency across the blending.

3 Data & Implementation

3.1 Asset Universe

The investable universe consists of eight liquid US-listed ETFs spanning major asset classes:

Table 1: Asset Universe

Ticker	Description	Class
SPY	S&P 500	Equity
QQQ	NASDAQ-100	Equity
EFA	MSCI EAFE	Intl Equity
VNQ	US REITs	Real Estate
TLT	20+ Year Treasury	Long Bond
IEF	7-10 Year Treasury	Mid Bond
LQD	Investment Grade Corp	Credit
GLD	Gold	Commodity

Daily adjusted closing prices are sourced from Yahoo Finance via **yfinance**. Data collection begins January 2005 and runs through December 2024, yielding 5,032 daily price observations and 5,031 daily log returns per asset. Log returns are computed as $r_t = \log(P_t/P_{t-1})$.

3.2 Evaluation Design

Data collection begins January 2005; the first 252 trading days serve as the burn-in period for rolling feature initialization (21-day volatility, 63-day correlation) and to satisfy the minimum HMM training window. Accordingly, the backtest period begins January 2006 and runs through December 2024 — 18 years. The sample is split into a training period (January 2006 – December 2018) and a held-out test period (January 2019 – December 2024).

Walk-forward retraining uses a quarterly frequency with an expanding training window. Of 78 theoretical quarterly retraining dates, 3 are excluded due to the 252-day minimum window, yielding 75 usable folds. The 2019–2024 slice is used as the held-out evaluation period.

3.3 Transaction Costs

A flat round-trip cost of 2 basis points is applied to all portfolio weight changes, including the cost of the initial allocation from an unfunded (cash) starting position:

$$r_t^{\text{net}} = r_t^{\text{gross}} - c \cdot \sum_i |w_{t,i} - w_{t-1,i}| \quad (17)$$

where $c = 0.0002$ and $w_{0,i} \equiv 0$ for all assets at the start of the simulated series. Average daily turnover is 3.79%, implying an annualized transaction cost of approximately $3.79\% \times 252 \times 0.0002 \approx 0.19\%$ per annum. This is conservative for liquid ETFs at institutional scale and should be treated as a lower bound for retail implementation.

3.4 Benchmarks

We compare against five benchmarks, all net of the same 2bp transaction cost model:

- **SPY**: Buy-and-hold S&P 500
- **Equal Weight**: Monthly-rebalanced equal allocation across all 8 assets
- **60/40**: 60% SPY / 40% IEF, monthly rebalanced
- **Momentum**: Monthly rebalancing to the top-3 assets by equal-weighted trailing 252-trading-day return, computed through the most recent month end (no skip-month adjustment is applied)
- **Risk Parity**: Static constrained risk parity over all 8 assets, using the same log-barrier formulation as the Bear regime optimizer

3.5 Risk-Free Rate

The time-varying risk-free rate is sourced from the 3-month US Treasury Bill yield ($\text{\textasciitilde IRX}$ via Yahoo Finance). It is used both in the maximum-Sharpe QP objective and in all reported Sharpe ratio calculations, ensuring consistency between optimization inputs and performance reporting.

3.6 Implementation

The full pipeline is implemented in Python using `hmmlearn` [hmmlearn developers, 2023] (v0.3.3) for HMM estimation with full covariance matrices, `cvxpy` [Diamond and Boyd, 2016] for convex optimization, `scikit-learn` [Pedregosa et al., 2011] for Ledoit-Wolf shrinkage, and `joblib` for parallel fold execution. All code is available at <https://github.com/AbdelkrimCode/regime-adaptive-portfolio>. Continuous integration via GitHub Actions runs 114 unit and integration tests on every commit, including coverage for cache-consistency checks, EM convergence diagnostics, and the risk-parity and bootstrap corrections discussed in Sections 2 and 5.

4 Results

4.1 Full-Period Performance (2006–2024)

Table 2 reports annualized performance metrics over the full backtest period. The regime-adaptive strategy achieves a Sharpe ratio of 0.561 versus 0.435 for SPY buy-and-hold. The primary outperformance is in drawdown control: maximum drawdown of -25.59% versus -59.58% for SPY, and a Calmar ratio of 0.215 versus 0.140.

Table 2: Full-Period Performance (2006–2024)

Metric	Ours	SPY	EW	60/40	Momentum	RP
Ann. Return	5.49%	8.34%	5.96%	6.88%	5.63%	4.81%
Ann. Vol	7.46%	19.51%	11.17%	11.17%	13.28%	7.56%
Sharpe	0.561	0.435	0.448	0.525	0.388	0.479
Max DD	-25.59%	-59.58%	-35.05%	-35.60%	-27.39%	-24.36%
Calmar	0.215	0.140	0.170	0.193	0.206	0.197

The strategy’s volatility (7.46% annualized) is dramatically lower than all equity-heavy benchmarks, reflecting the regime-conditional allocation to bonds and gold during Bear and Crash periods. The standalone Risk Parity benchmark now shows the best volatility and drawdown of any strategy in the table, including the main portfolio – a direct, verified consequence of a methodology correction described in Section 2: an earlier version of the Bear-regime optimizer was silently producing near-equal-weight allocations rather than actual risk parity, due to a position-limit cap being applied to the optimizer’s raw pre-normalization output rather than its final normalized weights.

4.2 Held-Out Test Period (2019–2024)

Table 3 reports performance on the held-out 2019–2024 period. The strategy’s Sharpe (0.261) trails SPY (0.677) by a wide margin. Part of this reflects the 2019–2024 period’s sustained equity bull market, unfavorable to any defensively-allocated strategy. Part of it is a direct, verified consequence of the risk-parity correction discussed in Section 2: the previous, buggy near-equal-weight Bear allocation happened to be more resilient to 2022’s simultaneous bond/equity selloff than the corrected, volatility-driven allocation is. Both effects are reported here in full, not just the more flattering one.

Table 3: Held-Out Performance (2019–2024)

Metric	Ours	SPY
Ann. Return	4.23%	14.87%
Ann. Vol	8.68%	19.91%
Sharpe	0.261	0.677
Max DD	-25.59%	-35.75%
Calmar	0.165	0.416

The maximum drawdown is identical in both the full period (Table 2) and the held-out period (Table 3). This is not a reporting error: the strategy’s peak-to-trough loss was realized during the 2020 COVID-19 crash, which falls within the 2019–2024 test window.

4.3 Frozen Out-of-Sample Evaluation

To rule out lookahead in the walk-forward procedure itself, we train a single HMM on data through December 2018, freeze all parameters, and apply the model statically to 2019–2024 without any retraining. Table 4 compares frozen versus walk-forward performance.

Table 4: Frozen vs Walk-Forward (2019–2024)

Metric	Walk-Forward	Frozen	SPY
Ann. Return	4.23%	5.09%	14.87%
Ann. Vol	8.68%	8.57%	19.91%
Sharpe	0.261	0.339	0.677
Max DD	-25.59%	-25.66%	-35.75%
Calmar	0.165	0.198	0.416

The frozen model outperforms walk-forward (0.339 vs 0.261 Sharpe), consistent with the full-period finding that walk-forward retraining does not inflate performance – if anything, it may be adapting to noise rather than signal in the test period.

4.4 Subperiod Analysis

Table 5 decomposes performance across three economically distinct subperiods.

Table 5: Subperiod Analysis

Period	Metric	Ours	SPY
GFC (2008–09)	Sharpe	0.408	-0.345
	Max DD	-10.90%	-56.38%
Bull (2013–19)	Sharpe	0.673	1.003
	Max DD	-7.67%	-19.82%
	Calmar	0.536	0.669
COVID+Rates (2020–24)	Sharpe	0.016	0.528
	Max DD	-25.59%	-35.75%
	Calmar	0.084	0.331

The GFC result is the most compelling and is essentially unaffected by the risk-parity correction: the portfolio gained +4.65% annualized versus SPY −15.49%, with maximum drawdown of −10.90% versus −56.38% for SPY. The 2008–2009 crisis was dominated by Crash/Bear regime days where a concentrated and a balanced risk-parity allocation behave similarly, so this result held up across both versions of the optimizer.

The COVID+rates subperiod is now the strategy’s weakest by a wide margin – Sharpe fell to 0.016 once the risk-parity correction was applied (from 0.231 under the earlier, buggy near-equal-weight optimizer). This is the clearest evidence of the tradeoff discussed in Section 2: the 2022 rate-driven bear market saw bonds and equities sell off simultaneously, and a volatility-driven risk-parity allocation has no mechanism to anticipate that kind of correlation regime change – it was, in this specific historical window, worse off for being methodologically correct.

4.5 Regime Distribution

Over the 2006–2024 backtest period (4,717 days with assigned regimes), the HMM assigns 1,261 days (26.7%) to Bull, 1,344 (28.5%) to Bear, 1,218 (25.8%) to Sideways, and 894 (19.0%) to Crash. Average regime duration ranges from 17.3 days (Bull) to 27.9 days (Crash). The

empirical transition matrix shows Bull-to-Sideways as the single most common transition (61.6% of exits from Bull), with recoveries from both Bear and Crash tending toward Bull (47.8% and 21.9% of exits, respectively) or Sideways rather than a direct reversal. Regime detection itself is unaffected by the risk-parity correction, since regime assignment happens upstream of portfolio construction; the counts here are consistent with prior reporting to within normal session-to-session data variability (Section 6).

We note that regime counts and downstream performance figures are not bit-reproducible across independent data downloads on different days: Yahoo Finance’s historical price history is occasionally revised retroactively, and because BIC-based state selection is a discrete choice, small price differences can occasionally shift which state count a fold selects. All figures in this paper come from a single, uninterrupted pipeline run for internal consistency.

5 Robustness Analysis

5.1 Bootstrap Significance Test

We test whether the strategy’s Sharpe ratio exceeds SPY’s using a paired block bootstrap [Politis and Romano, 1994] with the same block offsets applied to both portfolio and SPY returns, preserving cross-series correlation. Under 1,000 bootstrap replications with block length $b = 20$ trading days, the empirical p -value is 0.344 – the null hypothesis that the strategy’s Sharpe is no better than SPY’s cannot be rejected at conventional significance levels.

Sensitivity to block length: rerunning with $b = 10$, $b = 20$, and $b = 40$ produces p -values of 0.332, 0.344, and 0.347 respectively – stable across the range, confirming the result is not an artifact of block length choice.

This result is reported prominently because it is honest: the strategy does not statistically outperform SPY on a risk-adjusted basis. The value proposition is drawdown control and tail risk reduction, not Sharpe maximization. This p -value is computed by two independent implementations in the accompanying codebase (the main pipeline’s bootstrap module and a standalone stress-testing script); both currently agree to three decimal places, which we treat as a useful internal consistency check on the bootstrap procedure itself.

5.2 Walk-Forward Leakage Audit

We programmatically verify that no training fold observes future test data. For each of 75 quarterly folds, we record the last training date τ_{train} and the first test date τ_{test} . The condition $\tau_{\text{train}} < \tau_{\text{test}}$ holds for all 75 folds. Zero leakage is confirmed. The full audit is available at `results/walk_forward_audit.csv`.

5.3 Forward Filter vs Smoothed Posteriors

A controlled comparison holding the fitted HMM fixed at each fold and varying only the posterior computation method (causal forward filter vs. `hmmlearn`’s forward-backward smoothed posterior estimate) shows the causal filter trailing the smoothed variant by approximately 0.02 Sharpe – the expected direction of effect, since smoothed posteriors at time t have

access to observations $\mathbf{x}_{t+1}, \dots, \mathbf{x}_T$ within each fold’s test window and can therefore assign a better-fitting regime label in hindsight than any real-time system could compute.

A note on precision here: this specific diagnostic compares posterior computation methods and is independent of which downstream optimizer consumes those posteriors, so the qualitative direction and rough magnitude of this finding are expected to hold regardless of the risk-parity correction described in Section 2. However, the exact figures for this comparison were measured prior to that correction and have not been re-verified against the final, corrected pipeline. We report the qualitative finding with that caveat rather than restate a precise figure we have not re-confirmed.

5.4 Feature Ablation Study

Table 6 reports full pipeline results across four feature configurations, rerun under the final corrected pipeline. The baseline (21-day vol, 63-day correlation) is not the top performer on either metric in this configuration: `vol_10` outperforms baseline on both full-sample Sharpe (0.639 vs 0.561) and held-out Sharpe (0.333 vs 0.261). `vol_42` again shows the highest full-sample Sharpe of any variant (0.694) alongside a much weaker held-out Sharpe (0.150) – the same overfitting signature observed consistently across every version of this study regardless of which optimizer fix is in place. `skew_kurt` remains the weakest performer on both metrics.

Table 6: Feature Ablation Results

Feature Set	Full Sharpe	OOS Sharpe	Max DD
Baseline (<code>vol₂₁</code> , <code>corr₆₃</code>)	0.561	0.261	-25.59%
<code>vol₁₀</code> (<code>vol₁₀</code> , <code>corr₆₃</code>)	0.639	0.333	-22.58%
<code>vol₄₂</code> (<code>vol₄₂</code> , <code>corr₆₃</code>)	0.694	0.150	-25.98%
<code>skew+kurt</code>	0.495	0.098	-25.72%

That `vol_10` now outperforms the default configuration on both metrics is a legitimate finding worth flagging honestly: it is not, on its own, sufficient grounds to change the project’s default configuration, since a single sensitivity-sweep run is one realization of the session-to-session data variability discussed in Section 4, and switching defaults based on one such run risks exactly the kind of in-sample-driven decision this ablation study exists to guard against. We record this as a candidate worth investigating across multiple independent data pulls in future work, rather than adopting it immediately.

We additionally tested VIX and yield curve slope (5Y–3M Treasury spread) as supplementary macro features in an earlier iteration of this study; the qualitative finding was that these features improved full-sample fit but did not generalize out-of-sample. We have not re-verified the exact magnitude of that result under the current, corrected pipeline and report it here only qualitatively; the 3-feature baseline remains the default pending the `vol_10` investigation noted above.

5.5 Window Sensitivity Sweep

A 3×3 grid search over $\{10, 21, 42\}$ volatility windows and $\{42, 63, 126\}$ correlation windows, rerun under the final corrected pipeline, yields full-sample Sharpe in the range $[0.51, 0.70]$

and held-out Sharpe in $[0.08, 0.33]$. As noted above, the baseline configuration ($W_\sigma = 21$, $W_\rho = 63$) is not the top cell on either grid in this run; it remains within the broadly stable region of both grids rather than an outlier, which is the more relevant robustness property than any single cell’s raw score.

5.6 Stress Tests

Table 7: Stress Test Results

Metric	Portfolio	SPY
VaR (95%, daily)	-0.74%	-1.84%
CVaR (95%, daily)	-1.08%	-3.07%
Sharpe, $VIX > 30$	-0.496	-1.952
MC Sharpe mean	0.562	—
MC Sharpe std	0.228	—
MC 95% CI	$[0.139, 1.009]$	—

Daily VaR and CVaR at the 95th percentile remain less than half those of SPY, consistent with the drawdown control result. Sharpe during $VIX > 30$ periods also improved relative to earlier reporting under the pre-correction optimizer (-0.496 vs. a previously-reported -1.096) – both strategies still produce negative Sharpe during acute crisis periods, but the portfolio’s loss per unit of risk is now much closer to breakeven than before. This is a distinct effect from the COVID+rates subperiod weakness discussed in Section 4: VIX-spike episodes tend to be sharp and short, where a volatility-responsive risk-parity allocation reacts well, whereas the 2022 weakness was a grinding, multi-month correlation breakdown between bonds and equities that a volatility-only signal has no way to anticipate. Monte Carlo Sharpe (1,000 bootstrap paths, block-resampled to preserve autocorrelation) produces a mean of 0.562, closely matching the empirical estimate of 0.561, with no evidence of estimation bias.

5.7 PCA Regime Validation

Principal Component Analysis of the three features reveals that Bull, Bear, and Sideways regimes overlap heavily in the central region of PC1-PC2 space (39.3% and 33.3% variance explained respectively). Meaningful separation exists only in the left tail of PC1, corresponding to extreme negative returns and elevated volatility – Crash and severe Bear episodes. This analysis is unaffected by the risk-parity correction, since it operates on the HMM’s input features rather than on optimizer output.

This finding explains the non-significant bootstrap result: the HMM adds value primarily in tail events, not in distinguishing normal market conditions. Drawdown control is the correct framing for this strategy’s contribution.

6 Limitations

Gaussian emission assumption. Jarque-Bera tests reject normality for all four regimes ($p \approx 0$). Excess kurtosis is highest in the Crash regime (6.11) — clearly fat-tailed relative to

a Gaussian — with Bull, Bear, and Sideways closer to 1.8–2.3. A Student- t HMM [Zhang et al., 2004] would be more appropriate, particularly for the Crash regime, but is not natively supported by `hmmlearn`; implementation from scratch via EM with t -distribution updates is left for future work.

Markov assumption. HMMs assume the current regime depends only on the previous regime. Markets exhibit momentum, mean-reversion, and long-memory volatility clustering (GARCH effects) at multiple horizons, violating the first-order Markov property. A higher-order HMM or regime-duration model would capture these dynamics at the cost of significantly increased parameter count.

SPY circularity. All three features are derived from SPY returns. The HMM detects regimes using the same instrument it indirectly trades, introducing circularity in the signal. Macro features (VIX, yield curve slope) were tested as orthogonal alternatives but did not generalize out-of-sample. True signal orthogonality would require data from a different domain (e.g., economic survey indices, options market microstructure).

Expanding covariance window. All optimizers use an expanding window of returns through the retraining date, growing from approximately 3,260 to 4,780 days over the backtest. Early folds are dominated by 2008 volatility; late folds mix this with the much calmer post-2022 regime. A rolling window would be more stationary but introduces estimation instability in short samples. The expanding window is retained as the default; rolling-window sensitivity is available in the repository.

Crash regime assumes safe-haven decorrelation. The inverse-volatility crash optimizer allocates to IEF, TLT, and GLD under the assumption that these assets are negatively or uncorrelated with equities during crises. In 2022, all three sold off simultaneously with equities due to the rate shock, limiting protection. A min-CVaR optimizer over the safe-haven subset would be more robust to correlation shocks.

Transaction cost model. The flat 2bp cost model does not account for market impact, bid-ask spread dynamics, or securities lending costs. For daily rebalancing at retail scale, realized costs would be higher. Sensitivity analysis over cost assumptions is available in the accompanying repository.

Time period bias. The 2006–2024 sample includes the GFC, a prolonged low-volatility bull market, COVID-19, and an aggressive rate cycle. The strategy has not been evaluated in structurally different environments such as the 1970s stagflation or prolonged deflation. Out-of-sample evidence from such periods is unavailable by construction.

Statistical significance. The paired block bootstrap p -value of 0.344 confirms that excess Sharpe over SPY is not statistically significant. Any claim of outperformance should be framed around drawdown control — a claim supported by the data — rather than risk-adjusted returns.

Sample mean returns as optimizer input. The Maximum Sharpe optimizer uses sample mean returns, which have near-zero signal-to-noise ratio at daily frequency. Shrinkage toward a factor model or Black-Litterman prior would be more robust; this is left for future work.

Risk parity’s correction has a real cost, reported here rather than left out. The Bear regime optimizer (Section 2) now achieves close-to-equal risk contribution (0.0005–0.0014 percentage points of spread on real historical data, versus 14–23pp under an earlier, incompletely-fixed version). This is the most consequential finding of this audit round,

and it cuts both ways: the standalone Risk Parity benchmark improved on every metric once corrected, but the main strategy’s held-out Sharpe fell from 0.453 to 0.261 and the COVID+rates subperiod Sharpe fell to near-zero (0.016), because the previous, buggy near-equal-weight Bear allocation happened to be more resilient to 2022’s simultaneous bond/equity selloff than correctly-balanced risk parity is. This is not an implementation error in the corrected version. It is a real property of the correct methodology: risk parity allocates based on realized volatility and has no mechanism to anticipate a correlation regime change before it happens. We report this tradeoff in full rather than retain the more flattering, incorrect version of the optimizer.

Data is not bit-reproducible across download sessions. As discussed in Section 4, a fresh download of the same historical price series on a different day can shift regime counts and performance figures slightly, most likely due to Yahoo Finance backfilling minor historical corrections between sessions. Readers attempting to reproduce these results exactly should expect results representative of, rather than identical to, the figures reported here.

7 Conclusion

We presented a regime-adaptive portfolio framework built on a forward-filtered Hidden Markov Model and regime-conditional convex optimization. The central methodological contribution is the replacement of standard HMM smoothing with a strictly causal forward filter, eliminating the within-window lookahead bias present in most published HMM portfolio systems. Holding fitted models fixed, this correction costs approximately 0.02 Sharpe relative to forward-backward smoothed posteriors – the expected direction of effect, since smoothing has access to hindsight a live system would not.

Applied to eight liquid ETFs over 2006–2024, the strategy achieves a maximum drawdown of –25.6% versus –59.6% for SPY buy-and-hold – a 57% reduction in peak-to-trough loss. During the 2008–2009 Global Financial Crisis specifically, the portfolio gained +4.65% annualized versus SPY –15.49%, with maximum drawdown –10.9% versus –56.4% for SPY. A frozen out-of-sample evaluation confirms these results are not an artifact of walk-forward lookahead: a model trained exclusively on pre-2019 data achieves Sharpe 0.339 on the 2019–2024 period, ahead of the walk-forward estimate of 0.261 on the same period.

A significant part of this paper’s contribution is methodological transparency rather than performance alone. During the audit process described throughout this paper, a bug in the Bear-regime risk-parity optimizer was found, fixed, found to be incompletely fixed on a first pass, and fully corrected on a second – each stage verified against real historical data rather than assumed. The complete correction has a real cost: held-out Sharpe fell from 0.453 to 0.261, and the COVID+rates subperiod fell to near-zero (0.016), because the earlier, buggy near-equal-weight allocation happened to be more resilient to 2022’s correlated bond/equity selloff than correctly-balanced risk parity is. We report the corrected, less flattering numbers throughout this paper rather than the earlier, incorrect ones, and treat the full audit trail – including the intermediate, insufficient fix – as part of the paper’s contribution rather than something to leave out.

The strategy does not statistically outperform SPY on a risk-adjusted basis (bootstrap $p = 0.344$), and underperforms significantly on raw returns in the 2019–2024 period. PCA analysis of the feature space confirms that regime separation is weak in normal market

conditions and meaningful only in the left tail, explaining both the drawdown control result and the absence of Sharpe outperformance. The defensible claim of this work remains drawdown control, not return generation, and that claim held up – if anything, strengthened slightly – through every stage of this audit.

Future work includes replacing Gaussian emissions with Student- t distributions to better model fat-tailed returns, particularly for the Crash regime; investigating whether a shorter volatility window (Section 5 notes `vol_10` outperformed the default configuration in the most recent sensitivity sweep, though this alone is not sufficient grounds for a default change); incorporating a correlation-regime-aware signal into the Bear-regime optimizer specifically to address the 2022-style tail case now clearly documented; and extending the framework to higher-frequency data where order book dynamics may provide more orthogonal regime signals.

Acknowledgments

The author thanks the open-source maintainers of `hmmlearn`, `cvxpy`, and `yfinance` for the tools that made this implementation possible, as well as Mustapha Aziz Laroussi for inspiring this work.

References

- Andrew Ang and Geert Bekaert. Regime changes in asset allocation. *Journal of Financial Economics*, 64(3):347–393, 2002.
- Gérard Cornuéjols and Reha Tütüncü. *Optimization Methods in Finance*. Cambridge University Press, 2006.
- Steven Diamond and Stephen Boyd. Cvxpy: A python-embedded modeling language for convex optimization. *Journal of Machine Learning Research*, 17(83):1–5, 2016.
- Massimo Guidolin and Allan Timmermann. Asset allocation under multivariate regime switching. *Journal of Economic Dynamics and Control*, 31(11):3503–3544, 2007.
- James D Hamilton. A new approach to the economic analysis of nonstationary time series and the business cycle. *Econometrica*, 57(2):357–384, 1989.
- hmmlearn developers. hmmlearn 0.3.3: Hidden markov models in python. <https://hmmlearn.readthedocs.io>, 2023. Accessed 2024.
- Olivier Ledoit and Michael Wolf. A well-conditioned estimator for large-dimensional covariance matrices. *Journal of Multivariate Analysis*, 88(2):365–411, 2004.
- Fabian Pedregosa, Gaël Varoquaux, Alexandre Gramfort, et al. Scikit-learn: Machine learning in python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.
- Dimitris N Politis and Joseph P Romano. The stationary bootstrap. *Journal of the American Statistical Association*, 89(428):1303–1313, 1994.

- Florin Spinu. An algorithm for computing risk parity weights. *SSRN Working Paper*, 2013.
URL <https://ssrn.com/abstract=2297383>.
- Christopher M Turner, Richard Startz, and Charles R Nelson. A markov model of heteroskedasticity, risk, and learning in the stock market. *Journal of Financial Economics*, 25(1):3–22, 1989.
- Ming Zhang, Gagan Bhatt, and Thomas Berger. A new approach to the student- t hidden markov model. *Pattern Recognition Letters*, 25(12):1423–1432, 2004.